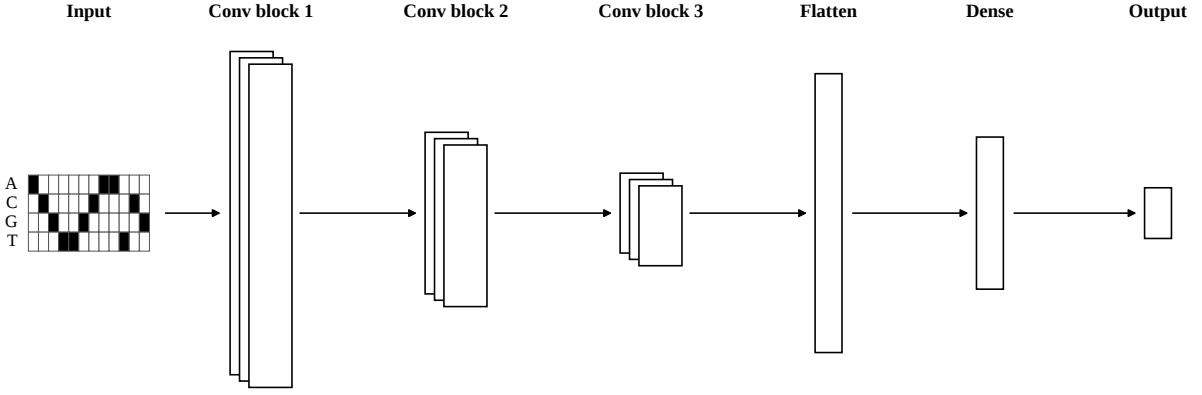

Loss: $L = \frac{1}{|\Omega|} \sum_{(i,a) \in \Omega} (\hat{z}_{ia} - z_{ia})^2$ over measured (sequence, assay) pairs. Adam, learning rate 10^{-3}, batch 256, early stopping (patience 5).